

```
[127]: import pandas as pd

train = pd.read_csv("/kaggle/input/train-and-test/CNNtrain.csv")
print(train.shape)
train.head()
```

(42000, 785)

```
[127... label pixel0 pixel1 pixel2 pixel3 pixel4 pixel5 pixel6 pixel7 pixel8 ... pixel774 pixel775 pixel776 pixel777 pixel778 pixel779 pixel780 pixel781 pixel782 pixel783
```

0	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	4	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows × 785 columns

```
[128]: # read test
test= pd.read_csv("/kaggle/input/train-and-test/CNNtest.csv")
print(test.shape)
test.head()
```

(28000, 784)

```
[128... pixel0 pixel1 pixel2 pixel3 pixel4 pixel5 pixel6 pixel7 pixel8 pixel9 ... pixel774 pixel775 pixel776 pixel777 pixel778 pixel779 pixel780 pixel781 pixel782 pixel783
```

0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	0	0	0	0	0

5 rows × 784 columns


```
[129]: # put labels into y_train variable
Y_train = train["label"]
# Drop 'label' column
X_train = train.drop(labels = ["label"],axis = 1)
```

```
[130]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(15,7))
g = sns.countplot(Y_train, palette="icefire")
plt.title("Number of digit classes")
Y_train.value_counts()
```


```
[130... label
1    4684
7    4401
3    4351
9    4188
2    4177
6    4137
0    4132
4    4072
8    4063
5    3795
Name: count, dtype: int64
```

Number of digit classes



```
:  
# plot some samples  
img = X_train.iloc[0].to_numpy()  
img = img.reshape((28,28))  
plt.imshow(img,cmap='gray')  
plt.title(train.iloc[0,0])  
plt.axis("off")  
plt.show()
```

1



```
:  
# Normalize the data  
X_train = X_train / 255.0  
test = test / 255.0  
print("x_train shape: ",X_train.shape)  
print("test shape: ",test.shape)
```

```
x_train shape: (42000, 784)  
test shape: (28000, 784)
```

```
:  
# Reshape  
X_train = X_train.values.reshape(-1,28,28,1)  
test = test.values.reshape(-1,28,28,1)  
print("x_train shape: ",X_train.shape)  
print("test shape: ",test.shape)
```

```
x_train shape: (42000, 28, 28, 1)  
test shape: (28000, 28, 28, 1)
```

```
:  
# Label Encoding  
from tensorflow.keras.utils import to_categorical # convert to one-hot-encoding  
Y_train = to_categorical(Y_train, num_classes = 10)
```

```
# Split the train and the validation set for the fitting  
from sklearn.model_selection import train_test_split  
X_train, X_val, Y_train, Y_val = train_test_split(X_train,  
                                                  Y_train,  
                                                  test_size = 0.1,  
                                                  random_state=2)  
  
print("x_train shape",X_train.shape)  
print("x_test shape",X_val.shape)  
print("y_train shape",Y_train.shape)  
print("y_test shape",Y_val.shape)
```

```
x_train shape (37800, 28, 28, 1)  
x_test shape (4200, 28, 28, 1)  
y_train shape (37800, 10)  
y_test shape (4200, 10)
```

```
from sklearn.metrics import confusion_matrix  
import itertools  
  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import Dense, Dropout  
from tensorflow.keras.layers import Flatten, Conv2D, MaxPool2D  
from tensorflow.keras.optimizers import RMSprop,Adam
```

```
model = Sequential()
# input layer
model.add(Conv2D(filters = 8, kernel_size = (7,7),padding = 'Same',
                  activation = 'relu', input_shape = (28,28,1)))
model.add(MaxPool2D(pool_size=(2,2)))
model.add(Dropout(0.25))
# hidden layer 1
model.add(Conv2D(filters = 16, kernel_size = (5,5),padding = 'Same',
                  activation = 'relu'))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Dropout(0.25))

# hidden layer 2
model.add(Conv2D(filters = 8, kernel_size = (3,3),padding = 'Same',
                  activation = 'relu'))
model.add(MaxPool2D(pool_size=(2,2), strides=(2,2)))
model.add(Dropout(0.25))

# fully connected
model.add(Flatten())
model.add(Dense(256, activation = "relu"))
model.add(Dropout(0.5))
model.add(Dense(10, activation = "softmax"))
```

```
# Define the optimizer
optimizer = Adam(learning_rate=0.001, beta_1=0.9, beta_2=0.999)
```

```
# Compile the model
model.compile(optimizer = optimizer ,
              loss = "categorical_crossentropy",
              metrics=["accuracy"])
```

```
model.summary()
```

Model: "sequential_10"

Layer (type)	Output Shape	Param #
conv2d_20 (Conv2D)	(None, 28, 28, 8)	400
max_pooling2d_20 (MaxPooling2D)	(None, 14, 14, 8)	0
dropout_31 (Dropout)	(None, 14, 14, 8)	0
conv2d_21 (Conv2D)	(None, 14, 14, 16)	3216
max_pooling2d_21 (MaxPooling2D)	(None, 7, 7, 16)	0
dropout_32 (Dropout)	(None, 7, 7, 16)	0
conv2d_22 (Conv2D)	(None, 7, 7, 8)	1160
max_pooling2d_22 (MaxPooling2D)	(None, 3, 3, 8)	0
dropout_33 (Dropout)	(None, 3, 3, 8)	0
flatten_11 (Flatten)	(None, 72)	0
dense_28 (Dense)	(None, 256)	18688
dropout_34 (Dropout)	(None, 256)	0
dense_29 (Dense)	(None, 10)	2570
Total params: 26034 (101.70 KB)		
Trainable params: 26034 (101.70 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
# data augmentation
datagen = ImageDataGenerator(
    featurewise_center=False, # set input mean to 0 over the dataset
    samplewise_center=False, # set each sample mean to 0
    featurewise_std_normalization=False, # divide inputs by std of the dataset
    samplewise_std_normalization=False, # divide each input by its std
    zca_whitening=False, # dimesion reduction
    rotation_range=5, # randomly rotate images in the range 5 degrees
    zoom_range = 0.1, # Randomly zoom image 10%
    width_shift_range=0.1, # randomly shift images horizontally 10%
    height_shift_range=0.1, # randomly shift images vertically 10%
    horizontal_flip=False, # randomly flip images
    vertical_flip=False) # randomly flip images

datagen.fit(X_train)
```



```
epochs = 15 # for better result increase the epochs
batch_size = 100
```


```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import ReduceLROnPlateau
```


```
# Fit the model
history = model.fit(datagen.flow(X_train,Y_train, batch_size=batch_size),
                    epochs = epochs,
                    validation_data = (X_val,Y_val),
                    steps_per_epoch=X_train.shape[0] // batch_size)
```

```
Epoch 1/15
378/378 [=====] - 26s 64ms/step - loss: 1.1463 - accuracy: 0.6057 - val_loss: 0.2646 - val_accuracy: 0.9331
Epoch 2/15
378/378 [=====] - 24s 64ms/step - loss: 0.5110 - accuracy: 0.8389 - val_loss: 0.1645 - val_accuracy: 0.9574
Epoch 3/15
378/378 [=====] - 24s 62ms/step - loss: 0.3919 - accuracy: 0.8759 - val_loss: 0.1245 - val_accuracy: 0.9643
Epoch 4/15
378/378 [=====] - 24s 64ms/step - loss: 0.3297 - accuracy: 0.8963 - val_loss: 0.1023 - val_accuracy: 0.9695
Epoch 5/15
378/378 [=====] - 24s 63ms/step - loss: 0.2986 - accuracy: 0.9070 - val_loss: 0.0892 - val_accuracy: 0.9726
Epoch 6/15
378/378 [=====] - 24s 63ms/step - loss: 0.2703 - accuracy: 0.9175 - val_loss: 0.0854 - val_accuracy: 0.9757
Epoch 7/15
378/378 [=====] - 24s 63ms/step - loss: 0.2474 - accuracy: 0.9243 - val_loss: 0.0747 - val_accuracy: 0.9764
Epoch 8/15
378/378 [=====] - 24s 62ms/step - loss: 0.2369 - accuracy: 0.9263 - val_loss: 0.0728 - val_accuracy: 0.9771
Epoch 9/15
378/378 [=====] - 24s 63ms/step - loss: 0.2266 - accuracy: 0.9310 - val_loss: 0.0734 - val_accuracy: 0.9767
Epoch 10/15
378/378 [=====] - 24s 63ms/step - loss: 0.2129 - accuracy: 0.9349 - val_loss: 0.0656 - val_accuracy: 0.9800
Epoch 11/15
378/378 [=====] - 24s 64ms/step - loss: 0.2095 - accuracy: 0.9367 - val_loss: 0.0633 - val_accuracy: 0.9800
Epoch 12/15
378/378 [=====] - 24s 62ms/step - loss: 0.2006 - accuracy: 0.9375 - val_loss: 0.0629 - val_accuracy: 0.9812
Epoch 13/15
378/378 [=====] - 24s 63ms/step - loss: 0.1954 - accuracy: 0.9420 - val_loss: 0.0617 - val_accuracy: 0.9814
Epoch 14/15
378/378 [=====] - 24s 63ms/step - loss: 0.1846 - accuracy: 0.9424 - val_loss: 0.0627 - val_accuracy: 0.9819
Epoch 15/15
378/378 [=====] - 23s 62ms/step - loss: 0.1862 - accuracy: 0.9439 - val_loss: 0.0606 - val_accuracy: 0.9810
```

```
# summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()

# summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```





```
# confusion matrix
import seaborn as sns
# Predict the values from the validation dataset
Y_pred = model.predict(X_val)
# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_pred,axis = 1)
# Convert validation observations to one hot vectors
Y_true = np.argmax(Y_val,axis = 1)
# compute the confusion matrix
confusion_mtx = confusion_matrix(Y_true, Y_pred_classes)
# plot the confusion matrix
f,ax = plt.subplots(figsize=(8, 8))
sns.heatmap(confusion_mtx, annot=True, linewidths=0.01,
            cmap="Greens",linecolor="gray", fmt= '.1f',ax=ax)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix")
plt.show()
```